

# Left join

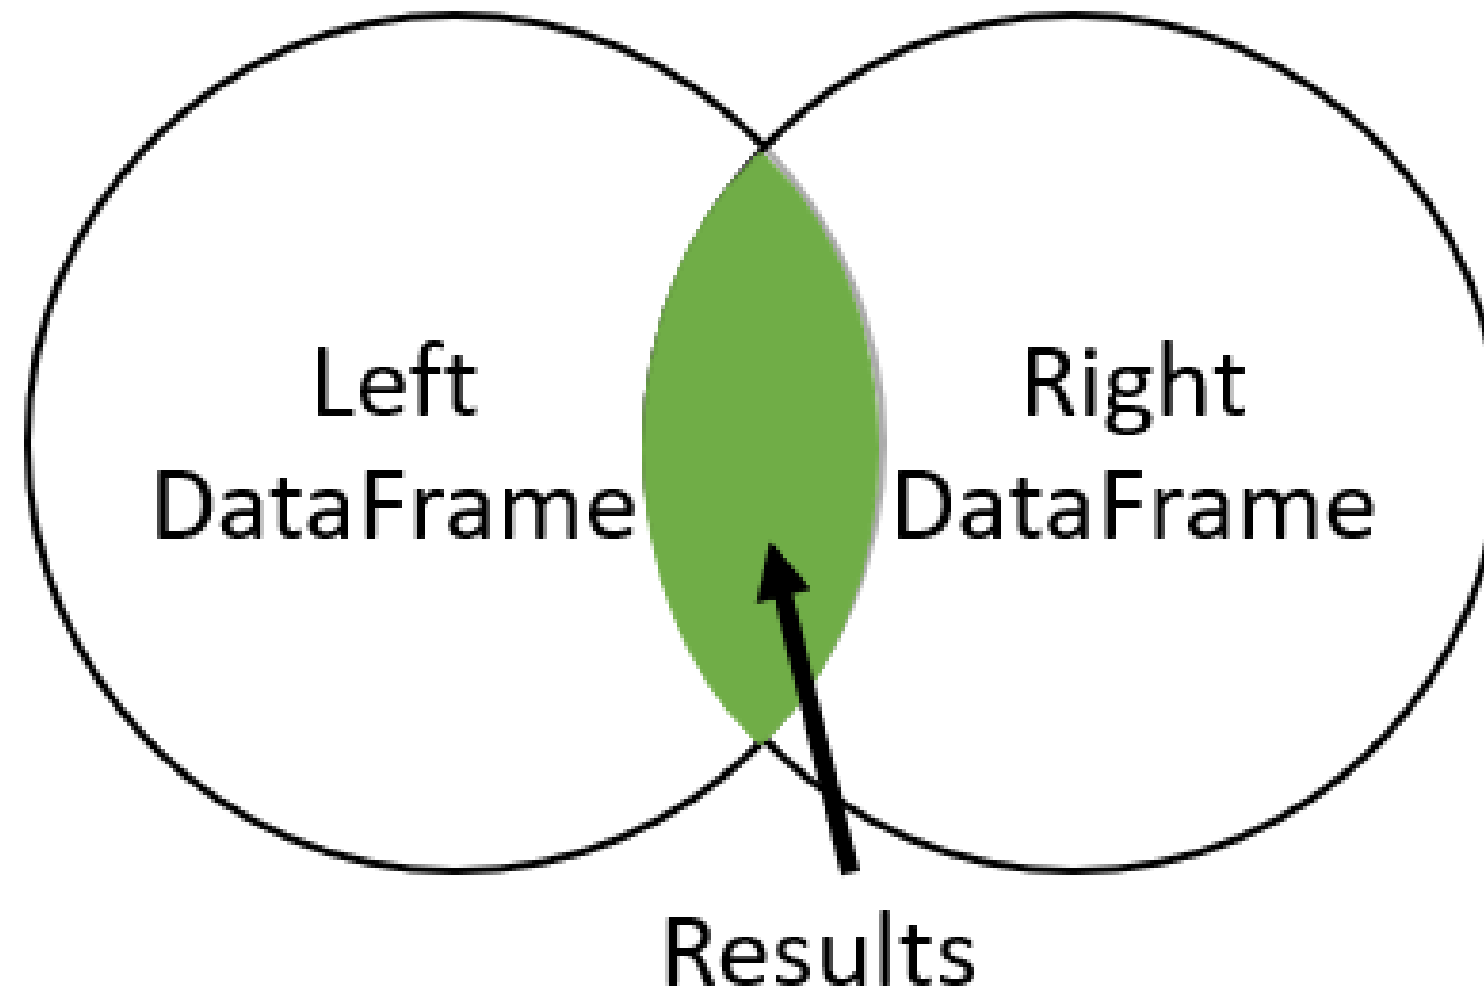
JOINING DATA WITH PANDAS



**Aaren Stubberfield**  
Instructor


## Quick review

# Inner Join



# Left join

## Left Join



A left join returns all rows of data from the left table and only those rows from the right table where key columns match.

# Left join

| Left Table |    |    | Right Table |    | Result Table |    |    |    |    |
|------------|----|----|-------------|----|--------------|----|----|----|----|
| A          | B  | C  | C           | D  |              | A  | B  | C  | D  |
| A2         | B2 | C2 | C1          | D1 | =            | A2 | B2 | C2 | D2 |
| A3         | B3 | C3 | C2          | D2 |              | A3 | B3 | C3 |    |
| A4         | B4 | C4 | C4          | D4 |              | A4 | B4 | C4 | D4 |
|            |    |    | C5          | D5 |              |    |    |    |    |

## 4. Left join

Here we have two tables named left and right. We want to use a left join to merge them on key column C. A left join returns all of the rows from the left table and only those rows from the right table where column C matches in both. Notice the second row of the merged table. The columns from the left table are filled in, while the column from the right table is not since there wasn't a match found for that row in the right table.

# New dataset

THE  
MOVIE  
DB



# Movies table

```
movies = pd.read_csv('tmdb_movies.csv')
print(movies.head())
print(movies.shape)
```

|   | id    | original_title  | popularity      | release_date |
|---|-------|-----------------|-----------------|--------------|
| 0 | 257   | Oliver Twist    | 20.415572       | 2005-09-23   |
| 1 | 14290 | Better Luck ... | 3.877036        | 2002-01-12   |
| 2 | 38365 | Grown Ups       | 38.864027       | 2010-06-24   |
| 3 | 9672  | Infamous        | 3.6808959999... | 2006-11-16   |
| 4 | 12819 | Alpha and Omega | 12.300789       | 2010-09-17   |

(4803, 4)

## 6. Movies table

Our first table, named `movies`, holds information about individual movies such as the title name and its popularity. Additionally, each movie is given an ID number. Our table starts with 4,803 rows of data.



# Tagline table

```
taglines = pd.read_csv('tmdb_taglines.csv')
print(taglines.head())
print(taglines.shape)
```

|   | id     | tagline  |
|---|--------|--|
| 0 | 19995  | Enter the World of Pandora.                    |
| 1 | 285    | At the end of the world, the adventure begins. |
| 2 | 206647 | A Plan No One Escapes                          |
| 3 | 49026  | The Legend Ends                                |
| 4 | 49529  | Lost in our world, found in another.           |

(3955, 2)

**7. Tagline table**  
Our second table is named **taglines**, which contains a movie ID number and the tag line for the movie. Notice that this table has almost 4,000 rows of data, so it contains fewer movies than the **movies** table.

# Merge with left join

```
movies_taglines = movies.merge(taglines, on='id', how='left')
print(movies_taglines.head())
```

|   | id    | original_title  | popularity      | release_date | tagline         |
|---|-------|-----------------|-----------------|--------------|-----------------|
| 0 | 257   | Oliver Twist    | 20.415572       | 2005-09-23   | NaN             |
| 1 | 14290 | Better Luck ... | 3.877036        | 2002-01-12   | Never undere... |
| 2 | 38365 | Grown Ups       | 38.864027       | 2010-06-24   | Boys will be... |
| 3 | 9672  | Infamous        | 3.6808959999... | 2006-11-16   | There's more... |
| 4 | 12819 | Alpha and Omega | 12.300789       | 2010-09-17   | A Pawsome 3D... |

## 8. Merge with left join

To merge these two tables with a left join, we use our merge method similar to what we learned in chapter 1. Here we list the movie table first and merge it to the taglines table on the ID column in both tables. However, notice an additional argument named 'how'. This argument defines how to merge the two tables. In this case, we use 'left' for a left join. The default value for how is 'inner', so we didn't need to specify this in Chapter 1 since we were only working with inner joins. The result of the merge shows a table with all of the rows from the movies table and a value for tag line where the ID column matches in both tables. **Wherever there isn't a matching ID in the taglines table, a null value is entered for the tag line. Remember that pandas uses NaN to denote missing data.**



# Number of rows returned

```
print(movies_taglines.shape)
```

```
(4805, 5)
```



## 9. Number of rows returned

After the merge, our resulting table has **4,805 rows**. This is because we are returning all of the rows of data from the movies table, and the relationship between the movies table and taglines is one-to-one. Therefore, in a one-to-one merge like this one, a left join will always return the same number rows as the left table.

Note: When performing a left join, the `.merge()` method returns a row full of null values for columns in the right table if the key column does not have a matching value in both tables.

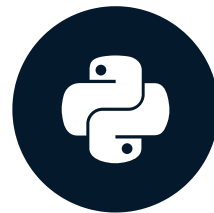
Note: A left join will return all of the rows from the left table. If those rows in the left table match multiple rows in the right table, then all of those rows will be returned. Therefore, the returned rows must be equal to if not greater than the left table. Knowing what to expect is useful in troubleshooting any suspicious merges.

# Let's practice!

JOINING DATA WITH PANDAS

# Other joins

## JOINING DATA WITH PANDAS




**Aaren Stubberfield**  
Instructor

# Right join

## 2. Right join

Right join will return all of the rows from the right table and includes only those rows from the left table that have matching values. It is the mirror opposite of the left join.



# Right join

Left Table

| A  | B  | C  |
|----|----|----|
| A2 | B2 | C2 |
| A3 | B3 | C3 |
| A4 | B4 | C4 |



Right Table

| C  | D  |
|----|----|
| C1 | D1 |
| C2 | D2 |
| C4 | D4 |
| C5 | D5 |

=

Result Table

| A  | B  | C  | D  |
|----|----|----|----|
|    |    | C1 | D1 |
| A2 | B2 | C2 | D2 |
| A4 | B4 | C4 | D4 |
|    |    | C5 | D5 |

# Looking at data

```
movie_to_genres = pd.read_csv('tmdb_movie_to_genres.csv')
tv_genre = movie_to_genres[movie_to_genres['genre'] == 'TV Movie']
print(tv_genre)
```

|       | movie_id | genre    |
|-------|----------|----------|
| 4998  | 10947    | TV Movie |
| 5994  | 13187    | TV Movie |
| 7443  | 22488    | TV Movie |
| 10061 | 78814    | TV Movie |
| 10790 | 153397   | TV Movie |
| 10835 | 158150   | TV Movie |
| 11096 | 205321   | TV Movie |
| 11282 | 231617   | TV Movie |

## 4. Looking at data

For this lesson, let's look at another table called `movie_to_genres`. Movies can have multiple genres, and this table lists different genres for each movie.



# Filtering the data

```
m = movie_to_genres['genre'] == 'TV Movie'
tv_genre = movie_to_genres[m]
print(tv_genre)
```

|       | movie_id | genre    |
|-------|----------|----------|
| 4998  | 10947    | TV Movie |
| 5994  | 13187    | TV Movie |
| 7443  | 22488    | TV Movie |
| 10061 | 78814    | TV Movie |
| 10790 | 153397   | TV Movie |
| 10835 | 158150   | TV Movie |
| 11096 | 205321   | TV Movie |
| 11282 | 231617   | TV Movie |

## 5. Filtering the data

For our right join example, let's take a sample of this data subsetting to develop a table of movies from the TV Movie genre.

# Data to merge

|   | id    | title           | popularity      | release_date |
|---|-------|-----------------|-----------------|--------------|
| 0 | 257   | Oliver Twist    | 20.415572       | 2005-09-23   |
| 1 | 14290 | Better Luck ... | 3.877036        | 2002-01-12   |
| 2 | 38365 | Grown Ups       | 38.864027       | 2010-06-24   |
| 3 | 9672  | Infamous        | 3.6808959999... | 2006-11-16   |
| 4 | 12819 | Alpha and Omega | 12.300789       | 2010-09-17   |

|       | movie_id | genre    |
|-------|----------|----------|
| 4998  | 10947    | TV Movie |
| 5994  | 13187    | TV Movie |
| 7443  | 22488    | TV Movie |
| 10061 | 78814    | TV Movie |
| 10790 | 153397   | TV Movie |

## 6. Data to merge

Our goal is to merge it with the movies table. We will set movies as our left table and merge it with the tv\_genre table. We want to use a right join to check that our movies table is not missing data. In addition to showing a right join, this example also allows us to look at another feature. Notice that the column with the movie ID number in the movies table is named `id`, and in the `tv_genre` table it is named `movie_id`. **The merge method has a feature to take this into account.**

# Merge with right join

```
tv_movies = movies.merge(tv_genre, how='right',  
                          left_on='id', right_on='movie_id')  
print(tv_movies.head())
```

|   | id     | title           | popularity | release_date | movie_id | genre    |
|---|--------|-----------------|------------|--------------|----------|----------|
| 0 | 153397 | Restless        | 0.812776   | 2012-12-07   | 153397   | TV Movie |
| 1 | 10947  | High School ... | 16.536374  | 2006-01-20   | 10947    | TV Movie |
| 2 | 231617 | Signed, Seal... | 1.444476   | 2013-10-13   | 231617   | TV Movie |
| 3 | 78814  | We Have Your... | 0.102003   | 2011-11-12   | 78814    | TV Movie |
| 4 | 158150 | How to Fall ... | 1.923514   | 2012-07-21   | 158150   | TV Movie |

## 7. Merge with right join


The code for this merge has some new elements. First of all, we set the `how` argument to `right` so that the merge performs a right join. Additionally, we introduce two new arguments, named `left_on` and `right_on`. They allow us to tell the merge which key columns from each table to merge the tables. We list `movies` as the left table, so we set `left_on` to `id` and `right_on` to `movie_id`. Our returned table has movies that match our table of `tv_genres`. There does not appear to be any null values in the columns from the `movies` table.

# Outer join

## 8. Outer join

Our last type of join is called an outer join. An outer join will return all of the rows from both tables regardless if there is a match between the tables.

One cool aspect of using an outer join is that, because it returns all rows from both merged tables and null where they do not match, you can use it to find rows that do not have a match in the other table.



# Outer join

Left Table

| A  | B  | C  |
|----|----|----|
| A2 | B2 | C2 |
| A3 | B3 | C3 |
| A4 | B4 | C4 |



Right Table

| C  | D  |
|----|----|
| C1 | D1 |
| C2 | D2 |
| C4 | D4 |
| C5 | D5 |



Result Table

| A  | B  | C  | D  |
|----|----|----|----|
|    |    | C1 | D1 |
| A2 | B2 | C2 | D2 |
| A3 | B3 | C3 |    |
| A4 | B4 | C4 | D4 |
|    |    | C5 | D5 |

## 9. Outer join

Here is a simple example of an outer join. Where the key column used to join the tables has no match, null values are returned. That is why in the result, the columns from the left table are missing in rows one and five, and in column D row three is missing.

# Datasets for outer join

```
m = movie_to_genres['genre'] == 'Family'  
family = movie_to_genres[m].head(3)
```

|   | movie_id | genre  |
|---|----------|--------|
| 0 | 12       | Family |
| 1 | 35       | Family |
| 2 | 105      | Family |

```
m = movie_to_genres['genre'] == 'Comedy'  
comedy = movie_to_genres[m].head(3)
```

|   | movie_id | genre  |
|---|----------|--------|
| 0 | 5        | Comedy |
| 1 | 13       | Comedy |
| 2 | 35       | Comedy |

## 10. Datasets for outer join

For an example of this, we filter the movie\_to\_genres table as before into two very small tables. One table has data on Family movies, and the other has Comedy movies.



# Merge with outer join

```
family_comedy = family.merge(comedy, on='movie_id', how='outer',  
                             suffixes=('_fam', '_com'))  
  
print(family_comedy)
```

|   | movie_id | genre_fam | genre_com |
|---|----------|-----------|-----------|
| 0 | 12       | Family    | NaN       |
| 1 | 35       | Family    | Comedy    |
| 2 | 105      | Family    | NaN       |
| 3 | 5        | NaN       | Comedy    |
| 4 | 13       | NaN       | Comedy    |

## 11. Merge with outer join

In this merge, we list the family table as the left table and merge it on the movie\_id column. The how argument is set to outer for an outer join. Both of our tables have the same column names. Therefore, we add suffixes to show what table the columns originated. In our result table, every row is returned for both tables and we see some null values. In our original comedy tables ID number 12 does not exist. Therefore a null is shown. Similarly, in our last row, movie ID 13 wasn't in the family dataset so it has a null.

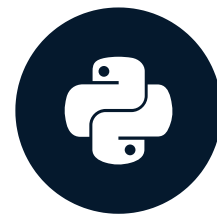
Note: Using an outer join, you were able to pick only those rows where the actor played in only one of the two movies.

# Let's practice!

JOINING DATA WITH PANDAS

# Merging a table to itself

JOINING DATA WITH PANDAS



**Aaren Stubberfield**  
Instructor

# Sequel movie data

```
print(sequel.head())
```

|   | id    | title        | sequel |
|---|-------|--------------|--------|
| 0 | 19995 | Avatar       | NaN    |
| 1 | 862   | Toy Story    | 863    |
| 2 | 863   | Toy Story 2  | 10193  |
| 3 | 597   | Titanic      | NaN    |
| 4 | 24428 | The Avengers | NaN    |



## 2. Sequel movie data

So when would you ever need to merge a table to itself? The table shown here is called `sequels` and has three columns. It contains a column for movie id, title, and sequel. The sequel number refers to the movie id that is a sequel to the original movie. For example, in the second row the movie is titled `Toy Story`, and has an id equal to 862. The sequel number of this row is 863. This is the movie id for `Toy Story 2`, the sequel to `Toy Story`. If we continue, 10193 is the movie id `Toy Story 3` which is the sequel for `Toy Story 2`.

# Merging a table to itself

| Left Table |             |        | Right Table |             |        | Result Table |             |          |       |             |          |
|------------|-------------|--------|-------------|-------------|--------|--------------|-------------|----------|-------|-------------|----------|
| id         | title       | sequel | id          | title       | sequel | id_x         | title_x     | sequel_x | id_y  | title_y     | sequel_y |
| 19995      | Avatar      |        | 19995       | Avatar      |        | 862          | Toy Story   | 863      | 863   | Toy Story 2 | 10193    |
| 862        | Toy Story   | 863    | 862         | Toy Story   | 863    | 863          | Toy Story 2 | 10193    | 10193 | Toy Story 3 |          |
| 863        | Toy Story 2 | 10193  | 863         | Toy Story 2 | 10193  |              |             |          |       |             |          |
| 597        | Titanic     |        | 597         | Titanic     |        |              |             |          |       |             |          |
| 24428      | The Ave...  |        | 24428       | The Ave...  |        |              |             |          |       |             |          |

Merge Columns

3. Merging a table to itself

If we would like to see a table with the movies and the corresponding sequel movie in one row of the table, we will need to merge the table to itself. In the left table, the sequel ID for Toy Story of 863 is matched with 863 in the ID column of the right table. Similarly, Toy Story 2 of the left table is matched with Toy Story 3 in the right table. We will talk more about this later, but the merge is an inner join. Therefore, we do not see Avatar and Titanic because they do not have sequels.



# Merging a table to itself

```
original_sequels = sequels.merge(sequels, left_on='sequel', right_on='id',  
                                suffixes=('_org', '_seq'))  
print(original_sequels.head())
```

|   | id_org | title_org       | sequel_org | id_seq | title_seq       | sequel_seq |
|---|--------|-----------------|------------|--------|-----------------|------------|
| 0 | 862    | Toy Story       | 863        | 863    | Toy Story 2     | 10193      |
| 1 | 863    | Toy Story 2     | 10193      | 10193  | Toy Story 3     | NaN        |
| 2 | 675    | Harry Potter... | 767        | 767    | Harry Potter... | NaN        |
| 3 | 121    | The Lord of ... | 122        | 122    | The Lord of ... | NaN        |
| 4 | 120    | The Lord of ... | 121        | 121    | The Lord of ... | 122        |

## 4. Merging a table to itself

We can think of it as merging two copies of the same table. All of the aspects we have reviewed regarding merging two tables still apply here. Therefore, we can merge the tables on different columns. We'll use the 'left\_on' and 'right\_on' attributes to match rows where the sequel's id matches the original movie's id. Finally, setting the suffixes argument in the merge method allows us to identify which columns describe the original movie and which describe the sequel. When we look at the results of the merge, the 'title\_org' and 'title\_seq' list the original and sequel movies, respectively. Here we listed the original movie and its sequel in one row.



# Continue format results

```
print(original_sequels[:, ['title_org', 'title_seq']].head())
```

|   | title_org       | title_seq       |
|---|-----------------|-----------------|
| 0 | Toy Story       | Toy Story 2     |
| 1 | Toy Story 2     | Toy Story 3     |
| 2 | Harry Potter... | Harry Potter... |
| 3 | The Lord of ... | The Lord of ... |
| 4 | The Lord of ... | The Lord of ... |

## 5. Continue format results

Now that we have our result table from the merge, we could select only the `title\_org`, and `title\_seq` columns, and we can see that we've successfully linked each movie to its sequel.

# Merging a table to itself with left join

```
original_sequels = sequels.merge(sequels, left_on='sequel', right_on='id',  
                                how='left', suffixes=('_org', '_seq'))  
print(original_sequels.head())
```

|   | id_org | title_org    | sequel_org | id_seq | title_seq   | sequel_seq |
|---|--------|--------------|------------|--------|-------------|------------|
| 0 | 19995  | Avatar       | NaN        | NaN    | NaN         | NaN        |
| 1 | 862    | Toy Story    | 863        | 863    | Toy Story 2 | 10193      |
| 2 | 863    | Toy Story 2  | 10193      | 10193  | Toy Story 3 | NaN        |
| 3 | 597    | Titanic      | NaN        | NaN    | NaN         | NaN        |
| 4 | 24428  | The Avengers | NaN        | NaN    | NaN         | NaN        |

## 6. Merging a table to itself with left join

... we can use the different types of joins we have already reviewed. Let's take the same merge from earlier but make it a left join. The 'how' argument is set in the merge method to left from the default 'inner'. Now the resulting table will show all of our original movie info. If the sequel movie exists in the table, it will fill out the rest of the row. If you compare this to our earlier merger, you now see movies like Avatar and Titanic in the result set.

# When to merge at table to itself

## Common situations:

- Hierarchical relationships
- Sequential relationships
- Graph data

### 7. When to merge at table to itself

You might need to merge a table to itself when working with tables that have a hierarchical relationship, like employee and manager.

You might use this on sequential relationships such as logistic movements.

Graph data, such as networks of friends, might also require this technique.

# Let's practice!

JOINING DATA WITH PANDAS

# Merging on indexes

JOINING DATA WITH PANDAS



**Aaren Stubberfield**  
Instructor

## 1. Merging on indexes

So far, we've only looked at merging two tables together using their columns. In this lesson, we'll discuss how to merge tables using their indexes. Often, the DataFrame indexes are given a unique id that we can use when merging two tables together.

# Table with an index

```
   id  title  popularity  release_date
0  257  Oliver Twist    20.415572   2005-09-23
1 14290  Better Luck ...    3.877036   2002-01-12
2 38365  Grown Ups      38.864027   2010-06-24
3  9672  Infamous       3.680896   2006-11-16
4 12819  Alpha and Omega  12.300789   2010-09-17
```

```
   title  popularity  release_date
id
257  Oliver Twist    20.415572   2005-09-23
14290  Better Luck ...    3.877036   2002-01-12
38365  Grown Ups      38.864027   2010-06-24
9672  Infamous       3.680896   2006-11-16
12819  Alpha and Omega  12.300789   2010-09-17
```



# Setting an index

```
movies = pd.read_csv('tmdb_movies.csv', index_col=['id'])  
print(movies.head())
```

|       | title           | popularity | release_date |
|-------|-----------------|------------|--------------|
| id    |                 |            |              |
| 257   | Oliver Twist    | 20.415572  | 2005-09-23   |
| 14290 | Better Luck ... | 3.877036   | 2002-01-12   |
| 38365 | Grown Ups       | 38.864027  | 2010-06-24   |
| 9672  | Infamous        | 3.680896   | 2006-11-16   |
| 12819 | Alpha and Omega | 12.300789  | 2010-09-17   |

## 3. Setting an index

There are different methods to set the index of a table, but if our data starts off in a CSV file, we can use the `index_col` argument of the `read_csv` method. This lesson will not focus on how to set a table index, but how to use that index to merge two tables together.

# Merge index datasets

|       | title           | popularity | release_date |
|-------|-----------------|------------|--------------|
| id    |                 |            |              |
| 257   | Oliver Twist    | 20.415572  | 2005-09-23   |
| 14290 | Better Luck ... | 3.877036   | 2002-01-12   |
| 38365 | Grown Ups       | 38.864027  | 2010-06-24   |
| 9672  | Infamous        | 3.680896   | 2006-11-16   |

|        | tagline         |
|--------|-----------------|
| id     |                 |
| 19995  | Enter the Wo... |
| 285    | At the end o... |
| 206647 | A Plan No On... |
| 49026  | The Legend Ends |

## 4. Merge index datasets

Recall our example to merge the movies and taglines tables using the id column with a left join. Let's recreate that merge using the index which is now the id for tables.

# Merging on index

Note: Merging on indexes is just like merging on columns, so if you need to merge based on indexes, there's no need to turn the indexes into columns first.

```
movies_taglines = movies.merge(taglines, on='id', how='left')
print(movies_taglines.head())
```

|       | title           | popularity | release_date | tagline         |
|-------|-----------------|------------|--------------|-----------------|
| id    |                 |            |              |                 |
| 257   | Oliver Twist    | 20.415572  | 2005-09-23   | NaN             |
| 14290 | Better Luck ... | 3.877036   | 2002-01-12   | Never undere... |
| 38365 | Grown Ups       | 38.864027  | 2010-06-24   | Boys will be... |
| 9672  | Infamous        | 3.680896   | 2006-11-16   | There's more... |
| 12819 | Alpha and Omega | 12.300789  | 2010-09-17   | A Pawsome 3D... |

## 5. Merging on index

Our merge statement looks identical to before. However, in this case we are inputting to the 'on' argument the index level name which is called 'id'. The merge method automatically adjusts to accept index names or column names. The returned table looks as before, except the 'id' is the index.

# Multindex datasets

```
samuel = pd.read_csv('samuel.csv',  
                    index_col=['movie_id',  
                              'cast_id'])  
  
print(samuel.head())
```

```
casts = pd.read_csv('casts.csv',  
                   index_col=['movie_id',  
                             'cast_id'])  
  
print(casts.head())
```

|          |         | name              |
|----------|---------|-------------------|
| movie_id | cast_id |                   |
| 184      | 3       | Samuel L. Jackson |
| 319      | 13      | Samuel L. Jackson |
| 326      | 2       | Samuel L. Jackson |
| 329      | 138     | Samuel L. Jackson |
| 393      | 21      | Samuel L. Jackson |

|          |         | character |
|----------|---------|-----------|
| movie_id | cast_id |           |
| 5        | 22      | Jezebel   |
|          | 23      | Diana     |
|          | 24      | Athena    |
|          | 25      | Elspeth   |
|          | 26      | Eva       |

## 6. Multindex datasets

Let's try a multindex merge. Here, we have two tables with a multindex that holds the movie ID and cast ID. The first table, named 'samuel', has the movie and cast ID for a group of movies that Samuel L. Jackson acted in. The second table, named cast, has the movie ID and cast ID for a number of movie characters. Let's merge these two tables on their multindex.

# Multindex merge

```
samuel_casts = samuel.merge(casts, on=['movie_id', 'cast_id'])  
print(samuel_casts.head())  
print(samuel_casts.shape)
```

|          |         | name              | character     |
|----------|---------|-------------------|---------------|
| movie_id | cast_id |                   |               |
| 184      | 3       | Samuel L. Jackson | Ordell Robbie |
| 319      | 13      | Samuel L. Jackson | Big Don       |
| 326      | 2       | Samuel L. Jackson | Neville Flynn |
| 329      | 138     | Samuel L. Jackson | Arnold        |
| 393      | 21      | Samuel L. Jackson | Rufus         |
| (67, 2)  |         |                   |               |

## 7. Multindex merge

In this merge, we pass in a list of index level names to the 'on' argument, just like we did when merging on multiple columns. Since this is an inner join, both the movie\_id and cast\_id must match in each table to be returned in the result. It's interesting to see that Samuel Jackson has acted in over 65 movies! That's a lot.

# Index merge with left\_on and right\_on

|       | title           | popularity | release_date |
|-------|-----------------|------------|--------------|
| id    |                 |            |              |
| 257   | Oliver Twist    | 20.415572  | 2005-09-23   |
| 14290 | Better Luck ... | 3.877036   | 2002-01-12   |
| 38365 | Grown Ups       | 38.864027  | 2010-06-24   |
| 9672  | Infamous        | 3.680896   | 2006-11-16   |

## 8. Index merge with left\_on and right\_on

There is one more thing regarding merging on indexes. If the index level names are different between the two tables that we want to merge, then we can use the `left_on` and `right_on` arguments of the merge method. Let's go back to our `movies` table, shown in the top panel, and merge it with our `movies_to_genres` table, shown in the lower panel.

|          | genre           |
|----------|-----------------|
| movie_id |                 |
| 5        | Crime           |
| 5        | Comedy          |
| 11       | Science Fiction |
| 11       | Action          |

# Index merge with left\_on and right\_on

```
movies_genres = movies.merge(movie_to_genres, left_on='id', left_index=True,  
                             right_on='movie_id', right_index=True)  
print(movies_genres.head())
```

|    | id | title      | popularity | release_date | genre           |
|----|----|------------|------------|--------------|-----------------|
| 5  | 5  | Four Rooms | 22.876230  | 1995-12-09   | Crime           |
| 5  | 5  | Four Rooms | 22.876230  | 1995-12-09   | Comedy          |
| 11 | 11 | Star Wars  | 126.393695 | 1977-05-25   | Science Fiction |
| 11 | 11 | Star Wars  | 126.393695 | 1977-05-25   | Action          |
| 11 | 11 | Star Wars  | 126.393695 | 1977-05-25   | Adventure       |

## 9. Index merge with left\_on and right\_on

In this merge, since we list the movies table as the left table, we set left\_on equal to id and right\_on equal to movie\_id. Additionally, since we are merging on indexes, we need to set left\_index and right\_index to True. These arguments take only True or False. Whenever we are using the left\_on or right\_on arguments with an index, we need to set the respective left\_index and right\_index arguments to True. The left\_index and right\_index tell the merge method to use the separate indexes.

# Let's practice!

JOINING DATA WITH PANDAS